

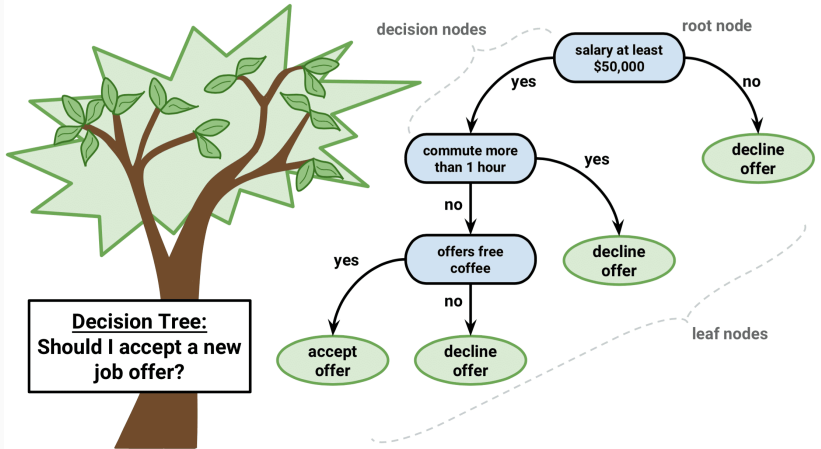
Statistical learning and Visualization

Boosting and Support Vector Machines

Maarten Cruyff

1. Boosting
2. Support Vector Machines

Classification tree



Classification trees

1. Recursive binary splitting algorithm
2. Splits features on basis of node *purity*




Figure 1: Sale of car seats (Yes/No)

Recursive binary splitting

Algorithm


1. Divide feature space in non-overlapping, rectangular regions
2. Choose splits that minimize *node impurity* (homogeneity of nodes)
3. Assign region to class with highest mode
4. Stop when node purity no longer increases

Algorithm is top-down and greedy, so

- high variance

Classification with trees or regression?

Depends on nature of relationship between classes and features



Package tree

Growing and plotting trees with function `tree()`

```
fit_tree <- tree(formula, data, split = c("deviance", "gini"))
```

```
plot(fit_tree)
```

```
text(fit_tree)
```

- minimization of deviance or gini impurity
- `text()` for adding labels to nodes

Methods to reduce variance:

1. Pruning

- cut branches with cross-validation and regularization

2. Bagging

- average predictions of bootstrapped trees

3. Random forests

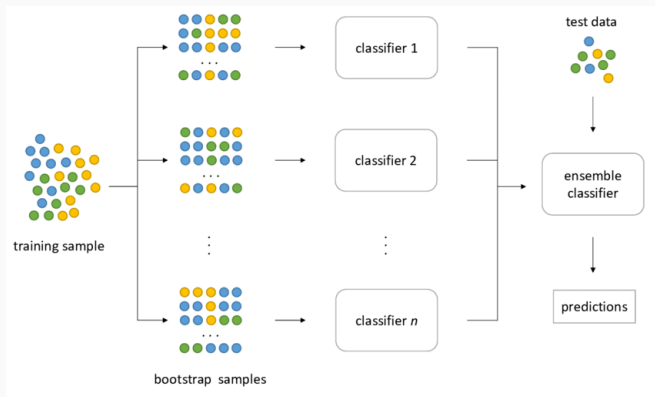
- average predictions of decorrelated bootstrapped trees

4. Boosting

- weighted combination of weak classifiers (small trees)

Random forest


1. Fit classification trees to B bootstrap samples
2. Average the predictions
3. Out-Of-Bag (OOB) as estimate validation error




Boosting

Algorithm

1. Apply a weak classifier (e.g. stump) to training data
2. Increase weights for incorrect classifications, and repeat
3. Classifier is linear combination of weak classifiers



Boosting vs bagging/random forest



Boosting with package gbm

Boosting a single model

```
boost <- gbm(formula, data, distribution)
```

```
predict(boost, newdata, type = c("link", "response"))
```

Boosting with package caret

Simple example

- distribution depends on response variable

```
gbm <- train(formula,  
              data,  
              distribution = "bernoulli",  
              method = "gbm")  
  
predict(gbm, newdata, type = "prob")
```

Support Vector Machines (SVM)

SVM for binary classification

Classifiers using support vectors

1. *maximal margin classifier*

- classes perfectly separable by hyperplane

2. *support vector classifier*

- allows for non-separable cases


3. *support vector machine*

- allows for non-linear boundaries

Hyperplane

Divides the feature space in two


- in two dimensions hyperplane is simply a line



Separating hyperplane

Perfectly separates the two classes of the outcome variable


- hyperplane not uniquely identified
- high variance



Maximal Margin Classifier

Identifies hyperplane by specification of a maximal margin


- points on margin are support vectors
- only works if cases are *separable*



Support Vector Classifier (SVC)

Allows for violations of the margin (soft margin)


- budget for violations is called *cost* (C)
- cases the wrong side of hyperplane contribute to the cost



Support Vector Machines (SVM)

Kernels allow for nonlinear hyperplanes, e.g.

- polynomial kernel (left)
- radial kernel (right)



SVM with package e1071

```
svm_train <- tune(svm, formula, data,  
                 degree = 3, #default  
                 coef0   = 0, #default  
                 cost    = 1, #default  
                 kernel  = c("linear", "polynomial", "radial"),  
                 ranges  = list(cost = <sequence>), etc.)
```

```
svm_train$best.model # performance summary
```

```
svm_class <- predict(svm_train, newdata, probability = TRUE)  
svm_prob  <- attr(svm_class, "probabilities")
```

- cost, degree and coef0 are tuning parameters
- ranges works similar as tuneGrid()

SVM classification plot

Compression hyperplane two dimensions

```
plot(svm_train$best.model, data, x1 ~ x2)
```

